


## ▼ sklearn in Python

Read in auto data

```
import pandas as pd
import sklearn
df = pd.read_csv("https://raw.githubusercontent.com/RyanBanafshay/Machine_Learning_Portfolio/main/ML%20with%20sklearn/Auto.csv")
df.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name	
0	18.0	8	307.0	130	3504	12.0	70.0	1	chevrolet chevelle malibu	
1	15.0	8	350.0	165	3693	11.5	70.0	1	buick skylark 320	
2	18.0	8	318.0	150	3436	11.0	70.0	1	plymouth satellite	
3	16.0	8	304.0	150	3433	12.0	70.0	1	amc rebel sst	
4	17.0	8	302.0	140	3449	NaN	70.0	1	ford torino	

Output the dimensions

```
df_ndim = df.shape
print("Number of dimensions: ", df_ndim)

Number of dimensions: (392, 9)
```

## ▼ Data Exploration

```
print("MPG:\n\n",df['mpg'].describe())
print("\nWeight:\n\n",df['weight'].describe())
print("\nYear:\n\n",df['year'].describe())
```

MPG:

```
count    392.000000
mean     23.445918
std       7.805007
min       9.000000
25%      17.000000
50%      22.750000
75%      29.000000
max      46.600000
Name: mpg, dtype: float64
```

Weight:

```
count    392.000000
mean    2977.584184
std     849.402560
min     1613.000000
25%     2225.250000
50%     2803.500000
75%     3614.750000
max     5140.000000
Name: weight, dtype: float64
```

Year:

```
count    390.000000
mean     76.010256
std       3.668093
min       70.000000
25%       73.000000
50%       76.000000
75%       79.000000
max       82.000000
Name: year, dtype: float64
```

Year has two less fields than MPG and Weight. MPG and Weight also have a relative small standard deviation. Year has a smaller range.

```
df.dtypes
```

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
```

```
df.cylinders = df.cylinders.astype('category').cat.codes
df = df.astype({"origin": 'category'})
df.dtypes
```

```
mpg          float64
cylinders     int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object
```

Drop NA values from the data

```
df=df.dropna()
df.isna().sum()
```

```
mpg          0
cylinders     0
displacement  0
horsepower    0
weight        0
acceleration  0
year          0
origin        0
name          0
dtype: int64
```

## ▼ Add Column

```
import numpy as np
avg = df.mpg.mean()
df['mpg_high'] = np.where(df.mpg > avg, 1, 0)
```

```
df = df.drop(columns=['name', 'mpg'])
print(df.head())
```

```
   cylinders  displacement  horsepower  weight  acceleration  year  origin  \
0          4         307.0         130   3504          12.0   70.0      1
1          4         350.0         165   3693          11.5   70.0      1
2          4         318.0         150   3436          11.0   70.0      1
3          4         304.0         150   3433          12.0   70.0      1
6          4         454.0         220   4354           9.0   70.0      1


   mpg_high
0         0
1         0
2         0
3         0
6         0
```

## ▼ Graphing the data

```
import seaborn as sb

sb.catplot(x = 'mpg_high', kind = 'count', data = df)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f7bfa5c6d90>
```



As we see from the visual representation, there seems to be very little difference in the data for mpg\_high.

```
sb.relplot(x = 'horsepower', y = 'weight', data = df, hue = df.mpg_high)
```


```
↳ <seaborn.axisgrid.FacetGrid at 0x7f7bfa48f3d0>
```



Horsepower and Weight much higher for the 0 mpg\_high. We also see that the weight typically corresponds with a higher horsepower.

```
sb.boxplot(x = 'mpg_high', y = 'weight', data = df)
```

<Axes: xlabel='mpg\_high', ylabel='weight'>



This boxplot shows us very similar results as the previous. There are a few outliers on mpg\_high 1

## ▼ Training the data

```
from sklearn.model_selection import train_test_split
```

```
X = df.iloc[:, 0:6]
```

```
y = df.iloc[:, 7]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=21, stratify=y)
```

```
print("Training size: ", X_train.shape)
```

```
print("Testing size: ", X_test.shape)
```

```
Training size: (272, 6)
```

```
Testing size: (117, 6)
```

## ▼ Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
logReg = LogisticRegression(solver = 'lbfgs', max_iter=100000)
```

```
logReg.fit(X_train, y_train)
```

```
logReg.score(X_train, y_train)
```

```
0.9007352941176471
```

## ▼ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
dTree = DecisionTreeClassifier()
```

```
dTree.fit(X_train, y_train)
```

```
▼ DecisionTreeClassifier
```

```
DecisionTreeClassifier()
```

```
predDT = dTree.predict(X_test)
```

```
accuracyDT = accuracy_score(y_test, predDT)
```

```
precisionDT = precision_score(y_test, predDT)
```

```
recallDT = recall_score(y_test, predDT)
```

```
f1DT = f1_score(y_test, predDT)
```

```
print("accuracy score: ", accuracyDT)
```

```
print("precision score: ", precisionDT)
```

```
print("recall score: ", recallDT)
```

```
print("f1 score: ", f1DT)
```

```
accuracy score: 0.9316239316239316
```

```
precision score: 0.9137931034482759
```

```
recall score: 0.9464285714285714
```

```
f1 score: 0.9298245614035087
```

## ▼ Neural Network

```
from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

from sklearn.neural_network import MLPClassifier

nn1 = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234)
nn1.fit(X_train_scaled, y_train)
```

```
▼                               MLPClassifier
MLPClassifier(hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234,
              solver='lbfgs')
```

```
prednn1 = nn1.predict(X_test_scaled)

from sklearn.metrics import confusion_matrix

accuracy_nn1 = accuracy_score(y_test, prednn1)
precision_nn1 = precision_score(y_test, prednn1)
recall_nn1 = recall_score(y_test, prednn1)
f1_nn1 = f1_score(y_test, prednn1)

print("accuracy score: ", accuracy_nn1)
print("precision score: ", precision_nn1)
print("recall score: ", recall_nn1)
print("f1 score: ", f1_nn1)

confusion_matrix(y_test, prednn1)
```

```
accuracy score:  0.9230769230769231
precision score:  0.9122807017543859
recall score:    0.9285714285714286
f1 score:        0.9203539823008849
array([[56,  5],
       [ 4, 52]])
```

## ▼ Second Neural Network

```
nn2 = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(4, 2), max_iter=500, random_state=1234)
nn2.fit(X_train_scaled, y_train)
```

```
▼                               MLPClassifier
MLPClassifier(hidden_layer_sizes=(4, 2), max_iter=500, random_state=1234,
              solver='lbfgs')
```

```
prednn2 = nn2.predict(X_test_scaled)

accuracy_nn2 = accuracy_score(y_test, prednn2)
precision_nn2 = precision_score(y_test, prednn2)
recall_nn2 = recall_score(y_test, prednn2)
f1_nn2 = f1_score(y_test, prednn2)

print("accuracy score: ", accuracy_nn2)
print("precision score: ", precision_nn2)
print("recall score: ", recall_nn2)
print("f1 score: ", f1_nn2)

confusion_matrix(y_test, prednn2)
```

```
accuracy score:  0.9145299145299145
precision score:  0.9107142857142857
recall score:    0.9107142857142857
f1 score:        0.9107142857142857
```

```
array([[56, 5],  
       [ 5, 51]])
```

In the second model, I lowered the number of hidden layers. This was ineffective at increasing the accuracy. Overall performance of the second model was generally worse than the first. I am assuming that the difference in both these models are minimal and not really reflective of how effective they are because the data size is so small.

## ▼ Analysis

The decision tree had the most accurate results, though not by such a significant amount that I would say it's the best algorithm. They all performed well with the data. A real test to see which would be better would be to use a variety of different and larger data sets and compare them using that.

If we compare the different metrics by class, we see:

Decision Tree:

- Accuracy = 0.9316239316239316
- Precision = 0.9137931034482759
- Recall = 0.9464285714285714

NN1:

- Accuracy = 0.9230769230769231
- Precision score = 0.9122807017543859
- Recall score = 0.9285714285714286

NN2:

- Accuracy = 0.9145299145299145
- Precision = 0.9107142857142857
- Recall = 0.9107142857142857

What algorithm to use is dependent on the data. I think a big reason why decision tree performed the best for this particular data is because decision trees are effective at handling non-linear datasets. Logistic regression had the worst accuracy, and in that case it needs a linear dataset to be fully effective. The neural network performed well but I think the main reason why it didn't perform as well is because NN's are more suited for larger more complex datasets.

Honestly, I prefer using sklearn in python over R. Python was designed to be easy to use for programmers used to more conventional programming languages such as Java or C++. Although R was convenient in how it was able to formulate results through internal functions, I never got used to the conventions of the language. Simple things like assigning variables felt more complicated than it needed to be. At a higher level, I would say both sklearn and R are very similar with no real preference one way or another. So the real defining reason why I prefer sklearn is simply that I find python a more familiar platform that has many similarities to Java and C++ with added conveniences.